



Contents lists available at ScienceDirect

The Journal of Systems and Software

journal homepage: www.elsevier.com/locate/jss

An event-driven high level model for the specification of laws in open multi-agent systems

Rodrigo Paes^{a,*}, Carlos Lucena^a, Gustavo Carvalho^a, Don Cowan^b

^aPUC-Rio, Rua M de S Vicente 225, Gávea 22453-900, Rio de Janeiro, Brazil

^bUniversity of Waterloo, 200 University Avenue West, Waterloo, Canada N2L 3G1

ARTICLE INFO

Article history:

Received 26 February 2007

Received in revised form 5 June 2008

Accepted 28 August 2008

Available online xxxx

Keywords:

Interaction laws

Multi-agent systems

Electronic institutions

Governance

Event-driven

Open systems

ABSTRACT

The agent development paradigm poses many challenges to software engineering researchers, particularly when the systems are distributed and open. They have little or no control over the actions that agents can perform. Laws are restrictions imposed by a control mechanism to deal with uncertainty and to promote open system dependability. In this paper, we present a high level event-driven conceptual model of laws. XMLaw is an alternative approach to specifying laws in open multi-agent systems that presents high level abstractions and a flexible underlying event-based model. Thus XMLaw allows for flexible composition of the elements from its conceptual model and is flexible enough to accept new elements.

© 2008 Published by Elsevier Inc.

1. Introduction

The agent development paradigm poses many challenges to software engineering researchers, particularly when the systems are distributed and open to accepting new modules that have been independently developed by third parties. Such systems have little or no control over the actions that agents can perform. As open distributed applications proliferate the need for dependable operation becomes essential.

There has been considerable research addressing the notion that the specification of such open multi-agent systems (MAS) should include laws that define behaviors in an open system (Carvalho et al., 2006b,c; Arcos et al., 2005). Laws are restrictions imposed by a control mechanism to tame uncertainty and to promote open system dependability (Minsky and Ungureanu, 2000; Paes et al., 2005). In this sense, such a mechanism should perform an active role in monitoring and verifying whether the behavior of agents is in conformance with the laws. We call this mechanism a governance mechanism. Examples of governance mechanisms are LGI (Minsky and Ungureanu, 2000), Islander (Esteva, 2003) and MLaw (Paes et al., 2006).

Governance for open multi-agent systems can be viewed as an approach that aims to establish and enforce some structure, set

of norms or conventions to articulate or restrain interactions in order to make agents more effective in attaining their goals or more predictable (Lindermann et al., 2006).

A governance approach has to deal with two important issues: a conceptual model (also called a domain language, or meta model) and the implementation mechanism that supports the specification and enforcement of laws based on the conceptual model. The content of this paper is mainly about the former.

In the conceptual model, the approach describes what elements designers can use when specifying the law. The model specifies the vocabulary and the grammar (or rules) that designers can use to design and implement the laws. The model has a decisive impact on how easy it is to specify and maintain the laws. It is the approach to design that largely determines the complexity of the related software. When the software becomes too complex, the software can no longer be well enough understood to be easily changed or extended. By contrast, a good design can make opportunities out of those complex features (Domain Language Inc, 2007).

In 1987, Minsky published the first ideas about laws (Minsky and Rozenshtein, 1987) and in 2000, he published a seminal paper about the role of interaction laws on distributed systems (Minsky and Ungureanu, 2000), which he called Law-Governed Interaction (LGI). Since then he has conducted further work and experimentation based on those ideas (Murata et al., 2003; Minsky, 2003, 2005). Although LGI can be used in a variety of application domains, its conceptual model is composed of abstractions basically

* Corresponding author. Tel.: +55 21 2540 6915x137.

E-mail addresses: rbp@inf.puc-rio.br (R. Paes), lucena@inf.puc-rio.br (C. Lucena), guga@inf.puc-rio.br (G. Carvalho), dcowan@csg.uwaterloo.ca (D. Cowan).

related to low level information about communication issues (such as the primitives disconnected, reconnected, forward, and sending or receiving of messages). While it can be possible to specify complex interaction rules based on such low level abstractions, they may not be adequate for the design of the laws pertaining to complex systems. This inadequacy occurs because, once the laws in the domain level are mapped to the many low level primitives, the original idea of the law is lost, as it is spread over many low level primitives. This is basically the problem of having a language that provides abstractions that are too far removed from the domain. When developing complex interactive systems, we need higher-level abstractions to represent laws in order to reduce complexity and achieve resultant productivity.

The Electronic institution (EI) (Esteva, 2003) is another approach that provides support for interaction laws. An EI has a set of high level abstractions that allow for the specification of laws using concepts such as agent roles, norms and scenes. Historically, the first ideas appeared when the authors analyzed the fish market domain (Noriega, 1997). They realized that to achieve a certain degree of regulation over the actions of the agents, real world institutions are needed to define a set of behavioral rules, a set of workers (or staff agents), and a set of observers (or governors) that monitor and enforce the rules. Based on these ideas, they proposed a set of abstractions and a software implementation. However, although EI provides high level abstractions, its model is quite inflexible with respect to change. The property of flexibility is quite important since the research in interaction laws is under constant evolution, and consequently the model that represents the law abstractions and their underlying implementation should also be able to evolve. One example of evolution is the use of laws for providing support to the implementation of dependability concerns. In this situation, the monitoring of laws may allow the detection of unexpected behaviors of the system and corresponding recovery actions.

The question is how can a conceptual model of laws evolve if we do not know in advance the nature of the changes? The answer to this question is that there is no way to foresee which parts are going to change. However, it is possible to use a basic underlying model that is inherently flexible. Event-based systems lead to flexible systems mainly because they avoid direct dependencies among the modules. Instead, the dependency is between the modules and the events they produce or consume.

In general, event-driven software design avoids any direct connection between the unit in charge of executing an operation and those in charge of deciding when to execute it. Event-driven techniques lead to a low coupling among modules and have gained acceptance because of their help in building flexible system designs (Meyer, 2003). In an event-based architecture, software components interact by generating and consuming events. When an event in a component (called source) occurs, all other components (called recipients) that have declared interest in the event are notified. This paradigm appears to support a flexible and effective interaction among highly reconfigurable software components (Cugola et al., 1998), and has been applied successfully in very different domains, such as graphical user interfaces, complex distributed systems (Cugola et al., 1998), component-based systems (Almeida et al., 2006) and software integration (Meier and Cahill, 2005). Many of these approaches use event-based systems to manage changes in the software that cannot be anticipated during design (Almeida et al., 2006; Batista and Rodriguez, 2000). Such changes are generally driven by a better understanding of the domain, and by external factors (such as strategic, political or budget decisions).

In this paper, we present a high level event-driven conceptual model of laws. The focus is to highlight the “high level” and “event-driven” aspects of the model, instead of presenting in detail the model itself. We do not claim that the abstractions of the proposed conceptual model are better than the ones in related ap-

proaches. Instead, we claim that the model is composed of a rich set of high level abstractions which enable, for instance, the specification of complex laws that can even interact with many current technologies (such as web services). The model is specified based on the event-driven paradigm. As a result, new elements can be easily introduced in the model.

The idea is that each element should be able to listen to and generate events. For example, if the model has the notion of norms, then this norm element should generate events that are potentially important to other elements, such as lifecycle events, norm activation, sanction applications and so on. The norm is also able to listen for events generated by other elements of the conceptual model, and then can react accordingly. For example, if in the conceptual model there is an element that models the notion of time, such as an alarm clock, then norms may listen to alarm clock notifications and their behavior becomes sensitive to time variations. This leads to very flexible and powerful relationships among the elements. Furthermore, if there is a need to introduce a new element in the model, then most of the work is restricted to connecting this new element to the events to which it needs to react, and to discover which events this new element should propagate.

The flexibility achieved by using the event-driven approach at a high level of abstraction is not present in the other high level approaches (Esteva, 2003; Dignum et al., 2004). The advantages claimed by the use of events as a modeling element are also present in Minsky’s approach (Minsky and Ungureanu, 2000) as a low level of abstraction. In this paper, we show in detail how to map our high level approach to Minsky’s in such a way that we illustrate that we can also achieve all the results he has produced so far (we are not addressing efficiency and security issues).

At the implementation level, we have developed middleware that supports the interpretation and enforcement of the specification and treats each element of the conceptual model as a component that is able to generate and sense events. The presentation of the middleware implementation is outside the scope of this paper and can be found in Paes et al. (2006, 2007).

This paper is organized as follows. Section 2 shows the proposed solution as a step in solving the stated problem. In Section 3, we relate our research to previous work, explaining how the problem of flexibility and evolution has been addressed. Section 4 shows two case studies where the model is applied and compared to related work. Finally, in Section 5, we present some discussions about the contents of this paper and future work.

2. XMLaw: an event-driven model

In this section, we present a partial view of XMLaw the conceptual model of laws and show how a new element can be added and connected to other elements through events. The full description of the XMLaw is described in XMLaw Specification (in press), Paes et al. (2005).

Fig. 1 shows the elements that compose the XMLaw conceptual model. The term conceptual model has the same meaning adopted by OMG¹ to refer to the UML conceptual model. This model can be viewed as composed of elements that cover many dimensions of the design of laws. As we could not find any related taxonomy in the literature, we are using the following ad-hoc classification:

Time – This dimension supports the specification of laws that are sensitive to time. For example, certain rules can have deadlines, expiration dates, and cyclical time-dependent behavior.

In the model, the Clock element represents this dimension. Clocks represent time restrictions or controls and can be used to activate other law elements. Clocks indicate that a certain period

¹ Object Management Group – <www.omg.org>.

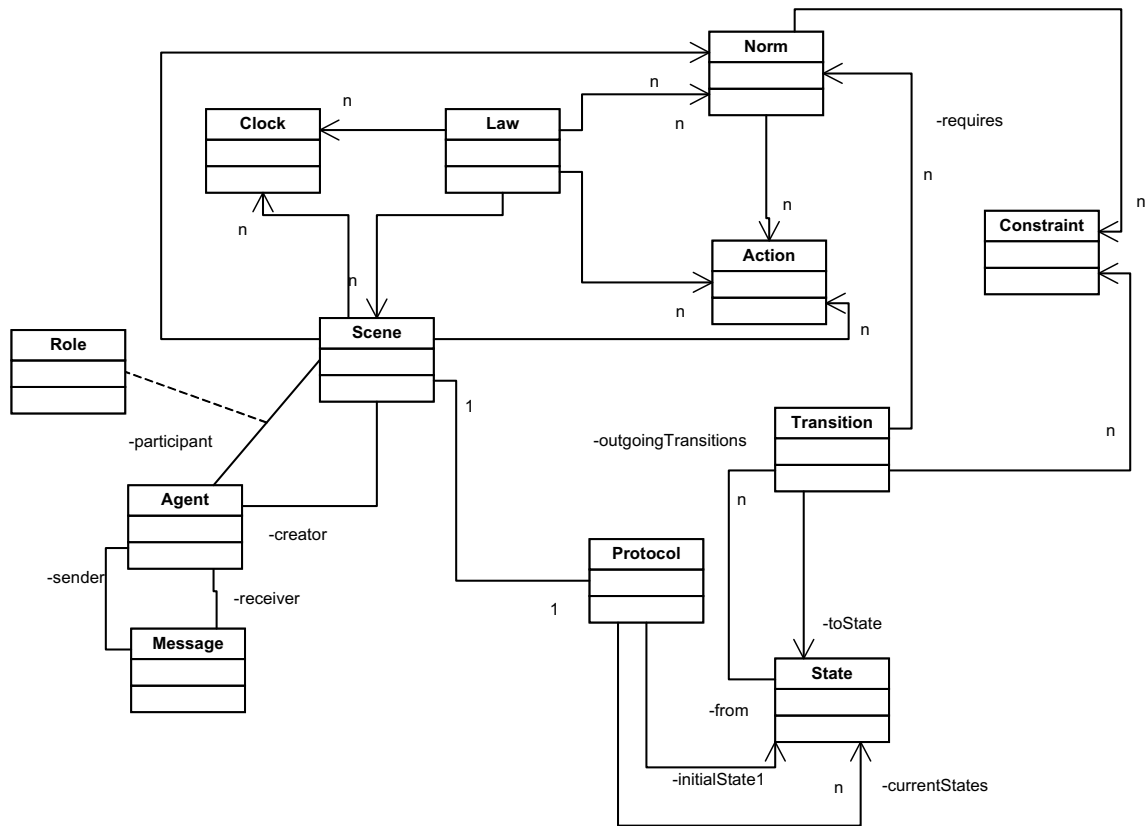


Fig. 1. The XMLaw conceptual model.

has elapsed producing *clock_tick* events. Once activated, a clock can generate *clock_tick* events. Clocks are activated and deactivated by law elements.

Social – This dimension supports the social relationships and interactions among the agents. Examples of social relationships are *master-slave* in some distributed systems and *employer-employee* in company environments.

The model has three elements in this dimension: Agent, Role and Message. The *agent* represents a software agent that is interacting with the other agents under the rules of the laws. There is no assumption about what language or architecture has been used to implement the agents. A role is a domain-specific representation of the responsibilities, abilities and expected behavior of an agent. It is useful to provide an abstraction for roles that is not related to the individuals playing the role. The Message element models a message exchanged among agents.

Structural – This dimension encompasses every type of structure used to describe the laws. They usually define modular contexts for which the laws are valid.

There are two elements in this dimension: Scene and Law. The idea of scenes is similar to the one in theater plays, where actors follow well defined scripts, and the whole play is composed of many sequentially connected scenes. A scene models a context of interaction where a protocol, actions, clocks and norms can be composed to represent complex normative situations. Furthermore, from the problem modeling point of view, a scene allows decomposition of the problem into smaller and more manageable pieces of information. The Law element is a context where all the other elements can be grouped. This is the most general element, and it does not belong to any other element.

Restrictive – This dimension contains elements with a focus on restricting the set of actions that agents are allowed to perform

in a given context. The model provides five elements in this dimension: Protocol, State, Transition, Norm, and Constraint.

A protocol defines the possible states through which an agent interaction can evolve. Transitions between states can be fired by any XMLaw event. Therefore, protocols specify the expected sequence of events in the interactions among the agents.

There are three types of norms in XMLaw: obligations, permissions and prohibitions. The obligation norm defines a commitment that software agents acquire while interacting with other entities. For instance, the winner of an auction is obliged to pay the committed value and this commitment might contain some penalties to avoid breaking this rule. The permission norm defines the rights of a software agent at a given moment. For example, the winner of an auction has permission to interact with a bank provider through a payment protocol. Finally, the prohibition norm defines forbidden actions for a software agent at a given moment; for instance, if an agent does not pay its debts, future participation in a scene will be prohibited.

Constraints are restrictions over norms or transitions and generally specify filters for events, constraining the allowed values for a specific attribute of an event. For instance, messages carry information that is enforced in various ways. A message pattern enforces the message structure fields. However, a message pattern does not describe what the allowed values for specific attributes are, but constraints can be used for this purpose. In this way, developers are free to build as complex a constraint as it is needed for their applications.

Service – This dimension is related to the interaction between the laws and the services that exist in the environment.

The action element belongs to this dimension. An action supports the definition of the moment when the mediator should call a domain-specific service.

Those elements of the model are connected through an event-based paradigm making it possible to achieve flexible behaviors through composition. For example, the norm element can be composed with the transition that creates an awareness of transition activation and thus behave properly. The pseudo code listed in Table 1 illustrates this idea. Other elements can also be composed to achieve complex behavior. For example, suppose that the buyer from the norm 1 in Table 1 now has to fulfill the obligation in at most 10 min. Then, this norm should incorporate some sense of time. In order to achieve that, we could compose the norm with the clock to reach the desired effect. Table 2 shows how this composition could be achieved. First, the clock listens to hear when the norm is given to the buyer. Then, the clock starts to count the ten minutes, and when this time has elapsed, the clock generates a *clock_tick* event. This *clock_tick* is then heard by norm2, which then prohibits any further interaction for the agent playing the buyer role.

In this scenario, the norm element was not originally designed or conceived to incorporate the notion of time, and the clock did not incorporate the notion of norms. The low coupling among these two elements caused by an event-based approach has lead to a flexible model of composition, even when the composition is not anticipated.

The evolution of the conceptual model of XMLaw has been influenced by the experiments we have been conducting. One special evolution applied to the original model was driven by the need for interacting with some services provided by the environment. In some cases, the laws could specify when and how to perform recovery actions; notify other stakeholders about changes in the law (through web services, for example), update database information and so on.

Based on these experiences, we have specified one more element, which is called Action. Actions are domain-specific Java code that runs in an integrated manner with XMLaw specifications. Actions can be used to plug services into a governance mechanism. For instance, a mechanism can call a debit service from a bank agent to charge the purchase of an item automatically during a negotiation. In this case, we specify in the XMLaw that there is a class that is able to perform the debit. Of course, this notion could also be extended to support other technologies instead of Java, such as direct invocation of web-services. Thus, these experiments demonstrate how a new action element can be included in the event-based model.

Within the *event-model*, the action can be integrated with the other elements by making actions able to listen to the other events. In this way, it would be possible to activate an action because of a clock activation, a norm activation, transitions and all the other events. On the other hand, to make the other elements able to react to actions, there is no need for changes, once the other elements can sense events; one can specify a listener for action activations.

Table 3 shows a situation where the action is activated by a transition, and then a clock is activated because of the action.

Although, the examples shown in this section are simple, they illustrate the consequences of having a flexible conceptual model. The conceptual model is usually mapped to some language (graph-

ical or textual) that allows the specification of laws. The second step is to build an interpreter that is able to read the specification and verify the compliance of the specification with the actual system behavior. The underlying event-based model can also be smoothly mapped to the implementation level, so the interpreter will also be flexible as it follows the same principles. We have implemented middleware, called M-Law (Paes et al., 2006) (Middleware for LAWs), that uses a component-based abstraction to represent each element of the conceptual model and an event-based model to make communication among the components possible.

2.1. The *event-driven* model definition

All the elements of the meta-model are able to sense and generate events, more precisely, let E be the set composed of the following elements: {Law, Scene, Norm, Clock, Protocol, State, Transition, Action, Constraint, Agent, Message, Role}, and let e denote an element of E , i.e., $e \in E$, then each is able to sense and generate events.

As can be seen in Fig. 2, every has the same basic lifecycle. It is important to notice that there is no restriction over which event can activate an element, this information provided through the law specification and therefore it is loosely coupled with the model.

Table 4 shows an example where the specification of the type of event that activates an element, in this case a clock, is just expressed in the law (for readability purposes the codes written in XMLaw presented in this paper use a simplified syntax which is more compact than the one used in early XMLaw publications). Line 16 says that a clock is activated (goes from idle to active state) when transitions t1 or t4 fire, and it is deactivated (goes from active to idle state) when transitions t2, t3 or t4 fire.

This very simple event mechanism has important consequences. It allows for a flexible and uncoupled composition of elements, and also allows for changes in the model. For example, when it is necessary to handle dependability concerns (de Gatti et al., 2006; Rodrigues et al., 2005).

3. Relating the model to other approaches

3.1. Relating the model to a lower level event-based approach

Minsky and Ungureanu (2000) proposed a coordination and control mechanism called law governed interaction (LGI). This mechanism is based on two basic principles: the local nature of the LGI laws and the decentralization of law enforcement. The local nature of LGI laws means that a law can regulate explicitly only local events at individual home agents, where a home agent is the agent being regulated by the laws; the ruling for an event e can depend only on e itself, and on the local home agent's context; and the ruling for an event can mandate only local operations to be carried out at the home agent. On the other hand, the decentralization of law enforcement is an architectural decision argued as necessary for achieving scalability.

Table 1
Pseudo code for activating a norm due to the firing of a transition

```
...
tl(s0,sl,message_arrival(ml))
...
norm1{
    give obligation to buyer when listen(transition_activation(tl))
}
...
```


Table 2

Pseudo code for composing the norm with the clock

```

...
tl(s0,s1,message_arrival(m1))
clock1{
  start to count when listen(norm_activation(norm1))
  count until 10 min and generate(clock_tick(clock1))
}
...
norm1{
  give obligation to buyer when listen(transition_activation(t1)) norm2{
    prohibit all interactions from buyer when listen(clock_tick(clock1))
  }
}
...

```


Table 3

Pseudo code for composing the norm with the clock

```

...
tl(s0,s1,message_arrival(m1))
clock1{
  start to count when listen(action_activation(action1))
  count until 5 s and generate(clock_tick(clock1))
}
action1{
  run ‘charge item in the bank’ when listen (transition_activation(t1))
}
...

```

**Fig. 2.** Generic law element lifecycle.

LGI has a rich set of events that can be monitored on each controller. Once these events are monitored, it is possible to use operations in order to implement the law. The union of events and operations is the conceptual model of LGI. The LGI conceptual model was conceived to deal with architectural decisions to achieve a high degree of robustness. This has lead to a model composed of low level primitives. Although the primitives are adequate for many classes of problems, it is necessary sometimes to use various primitives to achieve the desired effect. Once the laws become larger and more complex, it can be hard to maintain such a set of low level primitives.

One can think of LGI as a highly scalable virtual machine whose instructions are made of low level law elements. In this way, it would be possible, for example, to use high level abstractions of XMLaw to specify the laws, and in a second step map the specification to run on top of the LGI architecture. In order to illustrate this idea, we show how some of the elements that compose the conceptual model of XMLaw can be mapped to several of the LGI primitives. The illustration can easily be extended to cover all elements of the XMLaw model and the LGI architecture. We have summarized the main regulated events and operations of LGI in Tables 5 and 6.

Table 4

Defining the events that activate an element

```

...
08: t1{s1 -> s2, propose}
09: t2{s2 -> s3, accept}
10: t3{s2 -> s4, decline}
11: t4{s2 -> s2, propose}
...
16: clock{5000, regular, (t1,t4), (t2,t3,t4)}
...

```

Most of the events found in Table 5 are related to low level information about communication issues (*disconnected*, *reconnected*), sending or receiving of messages (*sent*, *arrived*), or state changes on the control state (*stateChanged*). From the point of view of the operations, they are also mostly concerned with low level instructions such as *forward*, *deliver*, *add* and so on. XMLaw has a rich set of high level abstractions. In Fig. 3 we show how to map some of abstractions to LGI instructions while preserving the meaning.

The protocol presented in Fig. 3 can be directly specified in XMLaw through the elements Protocol, State, Transition and Message. In order to achieve the same behavior in LGI, one can write the law as illustrated in Table 7. In this LGI law, we have introduced two terms: *currentState* and *event*. The term *currentState* models the current states of the protocol, and the term *event* simulates generation of events. For example, in the first line of Table 7, an agent *A* sends the message *m1* to the agent *B*; if the current state is *s0*, then state *s0* is removed from the list of current states, and state *s1* is added to the list. We also simulate the generation of a *transition_activation* event, and finally the message is forwarded.

Further examples in Table 8 show how some situations found in XMLaw can be mapped to the LGI approach, by incorporating the terms *event*, *currentState*, and *norm* in the LGI semantics expressed in prolog.

3.1.1. Global properties

According to Minsky (2005), “any policy that can be implemented via a central mediator – which can maintain the global interaction state of the entire community – can be implemented also via an LGI law”. As an example of a global property, suppose we encounter the situation in Fig. 4. In this example, 3 agents are interacting in the context of a specified protocol. Agent *A* sends the message *m1* to agent *B*. As agents *A* and *B* interact, their controllers update the current state. However, agent *C* has not participated of this interaction and therefore, its controller has not updated its current state. This causes an inconsistency between the agents *A* and *B* states and the state of agent *C*. This happens because the monitoring is performed in a decentralized way with no explicit synchronization.

The general way to overcome this problem is to have specific synchronization protocols such as the token ring in the Islander

Table 5
Main regulated events of LGI approach

Regulated events	
adopted	Represents the birth of an LGI agent – more specifically, this event represents the point in time when an actor adopts a given law L under which to operate thus becoming an L-agent
arrived	This event occurs when a message M sent by agent X to agent Y, arrives at the controller of Y. (The home of this event is agent Y – the receiver)
disconnected	This event occurs at the private controller of an agent when its actor has been disconnected
exception	This event may occur when the primitive operation, which has been invoked by the home agent, fails
obligationDue	This event is analogous to the sounding of an alarm clock, reminding the controller that a previously imposed obligation of a specified type is coming due. Obligations are imposed by means of the primitive operation imposeObligation
reconnected	This event occurs at the private controller of an agent when its previously disconnected actor has been reconnected
sent	This event occurs when the actor of x sends a message m addressed to an agent y operating under law L'. The sender x is the home of this event
stateChanged	This event occurs at an agent x when a pending state-obligation at x comes due
submitted	This event, which is a counterpart of the arrived event, occurs at an agent x, when an unregulated message m sent by some process at host h, using port p, arrives at x. It is, of course, up to law L under which x operates to determine the disposition of this message

This list is not intended to be complete.

Table 6
Main regulated operations of LGI approach

Operations	
Deliver	This operation, which has the form deliver ([x,L'],m,y), delivers to the home actor the message m, ostensibly sent by x, operating under law L'
Forward	Operation forward (x,m,[y,L']) sends the message m to Ty, the controller of the destination y – assumed here to operate under law L'; x is identified here as the ostensible sender of this message
Add	Adds term t to the CS
Remove	Removes from CS a term that matches t, if any. If there is no such term to be removed, this operation has no effect
Replace	Replaces a term t1 from CS, if any, with term t2. If there is no term t1 to be replaced, then this operation has no effect
Incr	Operation incr (f,d), locates a unary term f(n), and increments its argument by d
Decr	decr (f,d) is the implied counterpart of incr
ReplaceCS	Operation replaceCS (termList) replaces the whole control-state of the home agent with the specified list of terms
addCS	Operation addCS (termList) appends the terms in the list termList to the control state of the home agent
imposeObligation	imposeObligation (oType,dt,timeUnit) imposes an obligation of the specified type on the home agent, to come due after a delay dt, given in the specified time units
repealObligation	Operation repealObligation (oType) removes all pending obligations of type oType, along with all associated obligation-terms in DCS
imposeStateObligation	Operation imposeStateObligation (termList) would cause a stateChanged event to occur upon any change in any of the terms of the CS that are indicated by the termList parameter
repealStateObligation	Operation repealStateObligation (all) repeals the current state-obligation, as well as the corresponding audited (termList) term from the DCS

This list is not intended to be complete.




Fig. 3. Protocol example.

approach (Esteve, 2003). In LGI it could be achieved by introducing a central coordinator that receives all the messages, and therefore keeps a consistent global state. Table 9 outlines an implementation in LGI. The laws in XMLaw are specified from a global point of view (Paes et al., 2007). The approach presented in Table 9 has one advantage over the centralized mediator used in XMLaw, since the XMLaw mediator cannot protect itself against congestion be-

cause of some overactive participants which may lead to denial of service. Under LGI, on the other hand, the law may limit the frequency of messages that can be issued by any given participant. This limit can be locally enforced, and is less susceptible to congestion (Minsky and Ungureanu, 2000).

3.2. Relating the model to a not event-based approach

Electronic institutions (Esteve, 2003) are a technology to enforce and monitor the laws that apply to the agent society in a given environment. Several case studies were presented using this approach. They include a Fish Market system (Cuní et al., 2004), a Grid Computing Environment application (Ashri et al., 2006), and a Traffic Control application (Bou et al., 2006).

Table 7
Protocol specification in LGI

```

sent(A,m1,B) -> currentState(s0)@CS, do(remove(currentState(s0))),
do(add(currentState(s1))), do(add(event(t1,transition_activation))), do(forward)
sent(A,m2,B) -> currentState(s0)@CS,
do(remove(currentState(s0))), do(add(currentState(s2))),
do(add(event(t2,transition_activation))), do(forward)
sent(B,m3,A) -> currentState(s1)@CS, do(remove(currentState(s1))),
do(add(currentState(s2))), do(add(event(t3,transition_activation))), do(forward)
  
```

Table 8
XMLaw situation modeled using LGI

```

1. Upon the arrival of a message m1, a clock must be activated to fire an event in 5 s. The clock should be deactivated when the message m2 arrives
XMLaw
myXMLawClock{5000,regular,(m1),(m2)}
// Comments: myXMLawClock directly represents a clock. The clock abstraction already has
// the suitable structure to represent the time of activation, its type, and the messages
// that activate and deactivate it
LGI
arrived(X,m1,Y):- imposeObligation('myLGIclock",5)
arrived(X,m2,Y):- repealObligation("myLGIclock'')
// Comments: In LGI, it is necessary to combine different elements to simulate a clock
// in this case, it was necessary to combine the 'arrived' event with both
// imposeObligation and repealObligation to achieve the behavior of a clock
2. Fire transition t2 when the clock generates a clock_tick event. The transition changes the protocol from state s1 to state s2
XMLaw
t2{s1 -> s2, myXMLawClock}
// Comments: this scenario shows the composition between a clock and a transition. The
// activation of a clock can be directly linked to the activation of a transition
LGI
obligationDue('myLGIclock'):- currentState(s1)@CS, do(remove(currentState(s1))), do(add(currentState(s2))), do(add(event(t2,transition_activation))), do(forward).
// Comments: the obligationDue event happens when the time specified in
// the obligation expires. Note that for this relatively simple behavior it was necessary
// to use the obligationDue event, manipulate the control state (CS), create and
// manipulate the notion of currentState explicitly, generate the transition_activation
// event and forward the message to the addressee. All of these instructions had to be
// explicitly specified by the law designer. The main reason of this is although LGI is
// very flexible, its elements have a low level of abstraction, then to achieve higher
// level behavior it is necessary to combine many of these elements;
3. Declare a periodic clock that must generate an event every 5 s. This clock should be activated by the arrival of message m1 and deactivated by the arrival of message m2
XMLaw
myXMLawClock{5000,periodic,(m1),(m2)}
// Comments: Situation directly mapped to an XMLaw construct
LGI
arrived(X,m1,Y):- imposeObligation('myLGIclock",5)
arrived(X,m2,Y):- repealObligation('myLGIclock'')
obligationDue('myLGIclock'):- imposeObligation('myLGIclock",5)
// Comments: The obligationDue event happens when the time specified in
// the obligation expires. We can use a loop in the specification to
// simulate periodic clocks. In this example, first we 'declare' a clock, then
// when this clock expires, an obligationDue event is generated that in
// its turn activates another imposeObligation and so on
4. Message m1 activates transition t1. The transition t1 changes the protocol state from s1 to s2. The norm n1 must be activated when transition t1 is fired. The norm is given to the agent that
received the message m1.
XMLaw
t1{s1 -> s2,m1}
nl{$addressee,(t1)}
// Comments: In the first line, the message arrival is connected to the transition
// activation and in the second line, the transition activation is connected to the norm
// activation
LGI
sent(A,m1,Addressee) -> currentState(s1)@CS, do(remove(currentState(s1))), do(add(currentState(s2))), do(add(event(t1,transition_activation))), do(forward). imposeStateObligation(
event(t1,transition_activation)). stateChanged(event(t1,transition_activation)):- do(add(event(nl,norm_activation))), do(add(norm(nl,active,valid)))
// Comments: This example is a bit more complex than the first three examples. Then, it requires a combination of many elements of LGI in a relatively tricky manner to achieve the desired
behavior. The imposeStateObligation would cause an stateChanged event whenever the event(t1,transition_activation) term is added to the CS. Then, the third command states that when this
stateChanged happens, then the norm nl is made active

```

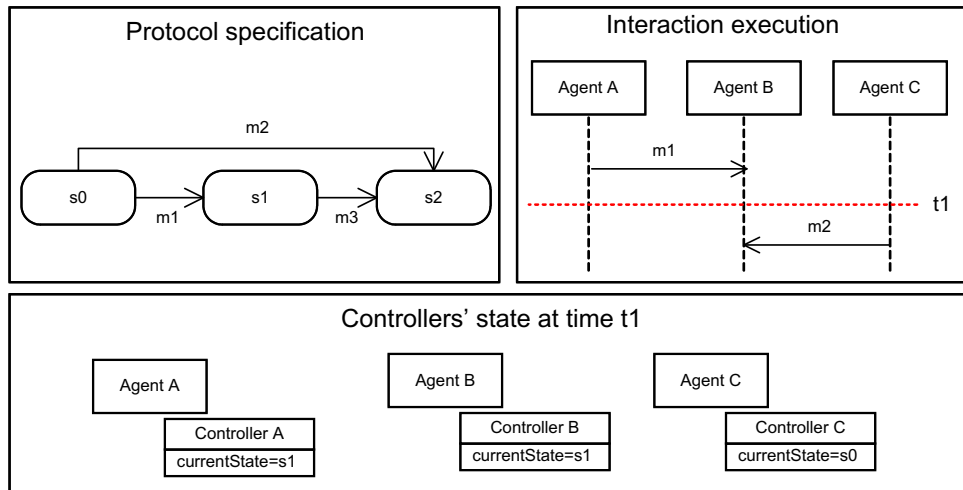


Fig. 4. Example of the need for synchronization of controllers to preserve global properties.

Table 9

Redirecting to a central coordinator

```
alias(coordinator,' law-coordinator@les.inf.puc-rio.br')
// Any message is forward to the central coordinator
sent(X,M,Y):- do(forward(X,[M,Y],#coordinator)), do(forward)
// Laws ... and redirection to the real addressee
arrived(#coordinator,A,m1,B) -> currentState(s1)@CS, do(remove(currentstate(s1))), do(add(currentState(s2))), do(deliver),
do(forward(A,m1,B))
```

Electronic institutions (EI) uses a set of concepts that have points of intersection with those used in XMLaw. For example, both EI scenes and protocol elements specify the interaction protocol using a global view of the interaction. The time aspect is represented in the Esteva's approach (Esteva, 2003) as timeouts. Timeouts allow activating transitions after a given number of time units have passed since a state was reached. On the other hand, because of the event-model, the clock element proposed in XMLaw can both activate and deactivate not only transitions, but also other clocks and norms. Connecting clocks to norms allows for a more expressive normative behavior; norms become time sensitive elements. Furthermore, XMLaw also includes the concept of actions, which allows execution of Java code in response to some interaction situations.

Table 10 compares the abstractions used in the conceptual model of both approaches. The goal of this comparison is to relate XMLaw better with an already existing well-known approach. The comparison shows that although they share a good set of concepts they have some important differences. For example, the notion of norms presented in EI (Garcia-Camino et al., 2005) is better defined than in XMLaw. On the other hand XMLaw has the concept of Actions that can be useful for making the laws behave more actively and integrated with services provided by the environment. However the major difference between the two models is the way that abstractions are related to compose the law. In EI there is a fixed set of relationships among the elements, and the way elements are used together is already defined in advance. A good example of this is the timeout abstraction. Timeout abstraction of EI is very similar to the clock abstraction of XMLaw. However, one can only use the timeout with transitions. If the underlying communication model among the elements were more flexible, it would be possible for example, to do the same thing as XMLaw and connect the timeout to the norm.

4. Case studies

In this section, we have chosen two examples published in the literature to illustrate the applicability of the XMLaw model. The

first example was already implemented and reported using the LGI approach. The second was also implemented and reported using the EI approach. By choosing these examples, we are able to compare the pros and cons of the various approaches directly.

4.1. Case study 1: *buyer team*

This case study was already implemented with LGI and presented in Minsky and Murata (2004). There are some modifications to the original problem description. We have eliminated the need for a certification authority. The example is described as follows.

"Consider a department store that deploys a team of agents, whose purpose is to supply the store with the merchandise it needs. The team consists of a manager, and a set of employees (or the software agents representing them) who are authorized as buyers and have access to a purchasing-budget provided to them.

Let us suppose that under normal circumstances, the proper operation of this buying team would be ensured if all its members comply with the following, informally stated, policy:"

1. The buying team is initially managed by a distinguished agent called *firstMgr*. But any manager of this team can appoint another agent authenticated as an employee as its successor, at any time, thus losing its own managerial powers.
2. A buyer is allowed to issue purchase orders (POs), taking the cost of each PO out of its own budget – which is thus reduced accordingly – provided that the budget is large enough. The copy of each PO issued must be sent to the current manager.
3. An employee can be assigned a budget by the manager, and can give some of that budget to other employees, recursively. In addition, the manager can reduce the budget of any employee *e*, as it sees fit, which freezes the budget of *e*, preventing others from increasing *e*'s budget. The budget of *e*, will only be able to increase again when the manager has changed.

Table 10
Relating EI and XMLaw conceptual models

Electronic institutions	XMLaw	Comments
Illocutory formulas	Message	They have different structures but they have the same semantics
EI vocabulary (ontology)	It is defined in the messages themselves, instead of separately	EI defines an explicit ontology of all the terms used in the conversation. XMLaw does not require this definition. In fact, the ontology in EI is the list of the types and functions expressed in Prolog or Lisp that the agents can use to communicate. It does not define the relationship and constraints among the concepts. In this way, the ontology in EI has a low contribution to the interoperability of the system
Internal roles	Not considered	Internal roles define a set of roles that will be played by staff agents that correspond to employees in traditional institutions. Since an EI delegates their services and duties to the internal roles, an external agent is never allowed to play any of the roles
External roles	Role	EI's external roles and XMLaw's role have the same goal. They are intended to represent the type of agents being mediated by the interaction laws
Relationships over roles	Not considered	In EI, relationships over roles can be specified, for instance, roles that cannot be played both at the same time. It is only possible because EI enforces the 'acquisition' of roles by agents. Although XMLaw allows role representation, it does not enforce this acquisition, which could be delegated to an authentication manager
Control over role playing	Control over role playing	Both approaches provide control over the minimum and maximum number of agents that can play a role in a scene
Scene	Scene	Both approaches have the notion of scene. In EI it is necessary to specify which agents are allowed either to enter or to leave a scene at some particular moment. In XMLaw, there is no need to specify the exit moments. This is because as agents can fail, or even exit at their own will, XMLaw considers exit moments as not necessary
Performative Structure	Not considered	The Performative structure is a special type of scene that accepts transitions from other scenes and has outgoing transitions to other scenes. They allow for a specification in which sequence scenes are expected to happen. In XMLaw, there is no such concept; however the notion of norm can be used to achieve similar a effect. Once in the end of a scene a norm could be activated and checked against the start of a new scene
Protocol	Protocol	In XMLaw the protocol is a non-deterministic finite state machine. In EI, the definition of the protocol does not clarify whether it is either deterministic or non-deterministic
State	State	States are quite similar, except that XMLaw has two types of final states: failure and success
Directed edge	Transition	The directed edge of EI (we used this name because EI has a transition element in the performative structure that has a different meaning) can be activated by illocution schemata, timeouts or constraints. In contrast, transitions in XMLaw can be activated by any event
Constraint	Constraint	In EI, constraints are specified as boolean expressions using an operator and two expressions: (op expr1 expr2). In XMLaw, they are implemented as domain-dependent Java Code
Time-out	Clock	Time-out allows provoking transitions after a given number of time units have passed since the state was reached. In contrast, clock is a general purpose clock that can also be used to provoke transitions to fire. But it can also be used, for example, to give an expiration period for a norm
Normative rules	Norms	Both approaches model notions of obligations, permissions and prohibitions. Normative rules model the notion of obligation by verifying: 'when an illocution is made and the illocution satisfies certain conditions THEN another illocution with other conditions must be satisfied in the future'. The norm in XMLaw can be used to prevent transition activations, actions activations and so on
Not considered	Actions	Actions can be used to plug services in the mediator. They can be activated by any event such as transition activation, norm activation and even action activation. The action specifies the Java class in charge of the functionality implementation
Not considered	Law	In XMLaw, the Law element is a global context where shared information among scenes and norms, clocks and actions can be used. In EI, a closed effect can be achieved through the performative structure

4.1.1. Messages

Item 1 of this policy is realized when the agent playing the manager role sends a message *transfer* to the employee that will be the successor. Item 2 happens when the buyer sends a *purchaseOrder* (*Amount*) message, where the *Amount* is the value of the purchase order that will be taken from the buyer's budget. Regarding item 3, an agent gives a budget to others by sending the message *giveBudget* (*Amount*). Then, the sender's budget will be reduced by *Amount* and the addressee's budget will be increased by *Amount*. Managers

can send the *removeBudget* (*Amount*) message. The effect of this message is to reduce by *Amount* the budget of the addressee.

4.1.2. XMLaw solution

XMLaw has abstractions to decompose the problem into small and more manageable pieces of information, and also to structure the steps of interactions of a complex conversational protocol. In this example, the interactions do not follow a pre-defined sequence, and the protocol is not too complex to justify decomposi-

tion into many small parts. We have specified the laws using one scene, which encapsulates the interaction protocol and a set of norms, actions and constraints. The complete specification can be found in Table 11, and the code for actions and constraints used in this specification can be found in Tables 12–15. We start the explanation by describing the general syntax and dynamics of the elements. We have also provided a graphical notation of the protocol based on UML statechart diagrams, which is shown in Fig. 5.

There are five types of messages that can be exchanged between the agents. These messages are specified in line 02 to line 06. The format of the message is *message-id{sender, receiver, content}*. The symbol * denotes any value. It is also possible to manipulate variables; variables are stored in the context. Each scene has its own context, and there is also a general law context. They form a hierarchy of contexts.

The message specified in line 02 means that the sender will be assigned to the *budgetOwner* variable, the receiver can be any agent, and the content has the form *purchaseOrder (value)*, where the value will be assigned to the variable *amount*. The messages are used to activate the five transitions of the protocol (lines 09 to 13). However, transitions t1, t2, t3, and t4 have constraints that will be checked before they fire and will only fire if the constraints are satisfied. The constraints are specified in lines 14 and 15. Once the transitions fire, some of them activate actions, as can be seen from line 16 to 18. Finally, transition t2 also needs the norm specified in line 19 in order to fire.

Transition t1 controls the purchase orders stated in item 02 of the policy. In order to be activated, the constraint *enoughMoney* referred to in line 09 verifies if the sender identified by the variable *budgetOwner* (line 02) has enough money to issue the order (Table 14). Transition t2 controls the item 03 of the policy. It refers to the situation where an employee or a manager gives a budget to another employee. Then, t2 first verifies if the sender has enough money to transfer (Table 14), then checks to see if the employee that is about to receive the money is allowed to receive money as specified in policy 03. This verification is done through the norm *increaseBudgetProhibition* (line 19). This norm is given to an employee when the manager sends a *removeBudget* message and this message activates the transition t3. In XMLaw it can be seen in line 19, where t3 is the transition that activates the norm, and t4 is the transition that deactivates it. Therefore, if the norm is active, the transition t2 is not fired. If the agent has not such a norm, then the transition t2 will fire. Once t2 is fired, action *changeBudget* is activated (line 18). The code of this action can be found in Table

Table 12

Constraint that verifies if the agent is in fact the current manager

```
class CheckTransfer implements IConstraint{
    public boolean constrain(ReadOnlyContext ctx){
        String actualManager = ctx.get('actualManager');
        String currentMgr = ctx.get('manager');
        if (! actualManager.equals(currentMgr)){
            return true;// constrains, transition should not fire
        }
        return false;
    }
}
```

Table 13

Action that switches the current manager to the employee

```
class SwitchManager implements IAction{
    public void execute(Context ctx){
        String employee = ctx.get('employee');
        ctx.put('actualManager', employee);
    }
}
```

15. Transition t4 is activated when the manager sends a *transfer* message to an employee. In the protocol, the constraint *checkTransfer* (Table 12) guarantees that the sender of the message is in fact the manager. If the transition t4 fires, then the action *switchManager* (Table 13) is executed, and the norm *increaseBudgetProhibition* is deactivated. The *switchManager* action updates the current manager, and once the *increaseBudgetProhibition* is not active, employees that have gained this norm, now are free again to receive budget through the message *giveBudget*.

4.1.3. Discussion

The most important part of this case study is Table 11. It is this table that contains the elements of the XMLaw conceptual model. The code in this table is mostly declarative and is concerned with high level abstractions such as interaction protocol, actions, constraints and norms. It was possible to express the rules in twenty lines of instructions, which are relatively simple to understand even for those not well versed in the XMLaw language. Even with the Java code needed to implement the actions and constraints, most of the time the designer can focus on the law specification of Table 11 and use the actions and constraints as components to achieve the desired functionality. The event-model of communication is present in most of the declarations. For example, line 19

Table 11
XMLaw Code

```
01: generalScene{
02:   PO{*budgetOwner, *purchaseOrder($amount)}
03:   removeBudget{manager,$budgetOwner,removeBudget($amount)}
04:   giveBudget{*budgetOwner,$receiver,giveBudget($amount)}
05:   transfer{*manager,$employee,transfer}
06:   end{*sender,$receiver,end}
07:   s1{initial}
08:   s2{success}
09:   t1{s1 -> s1, PO, [enoughMoney]}
10:   t2{s1 -> s1, giveBudget, [enoughMoney],[increaseBudgetProhibition]}
11:   t3{s1 -> s1, removeBudget, [enoughMoney]}
12:   t4{s1 -> s1, transfer, [checkTransfer]}
13:   t5{s1 -> s2, end}
14:   enoughMoney{br.pucurio.EnoughMoney}
15:   checkTransfer{br.pucurio.CheckTransfer}
16:   forwardMessage{(t1), br.pucurio.ForwardMessage}
17:   switchManager{(t4), br.pucurio.SwitchManager}
18:   changeBudget{(t2,t3), br.pucurio.ChangeBudget}
19:   increaseBudgetProhibition{*budgetOwner, (t3),(t4)}
20:}
```

Table 14

Constraint that verifies if the one who is giving money has enough money to give

```

class EnoughMoney implements IConstraint{
    public boolean constrain(ReadonlyContext ctx){
        String budgetOwner = ctx.get('budgetOwner');
        double currentBudget = Double.parseDouble(ctx.get(budgetOwner));
        double amount = ctx.get('amount');
        double diff = currentBudget - amount;
        if(diff<0){
            // constrains, not enough money. Transition should not fire
            return true;
        }
    }
}

```

Table 15

Action that updates budgets both for giveBudget and removeBudget

```

class ChangeBudget implements IAction{
    public void execute(Context ctx){
        String budgetOwner = ctx.get('budgetOwner');
        double currentBudget = Double.parseDouble(ctx.get(budgetOwner));
        double amount = ctx.get('amount');
        // update the owner's budget (the one who is given money)
        budget.put(budgetOwner, currentBudget - amount);
        String receiver = ctx.get('receiver');
        if (receiver!=null){// if there is a receiver
            double receiverBudget = Double.parseDouble(ctx.get(receiver));
            // update the receiver's budget
            ctx.put(receiver, receiverBudget+amount);
        }
    }
}

```

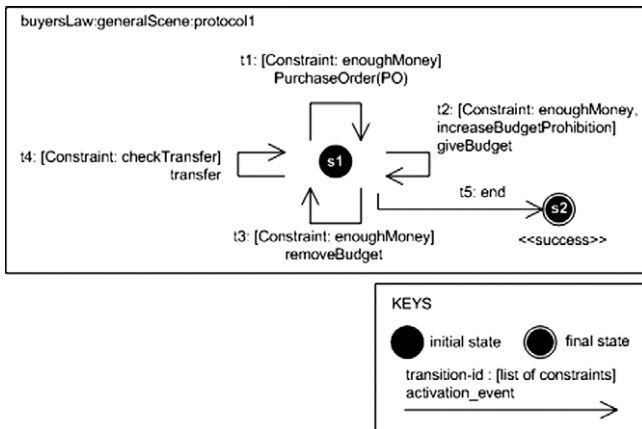


Fig. 5. Interaction protocol.

ders... provided that the budget is large enough" is directly mapped to the XMLaw Constraint element *enoughMoney* used in line 09.

4.2. Case study 2: *conference center*

This case study was implemented with EI and presented in Esteva (2003), Rodriguez-Aguilar et al. (2003). The example is described as follows:

"A conference takes place in a physical setting, the conference centre, where different activities take place in different locations by people that adopt different roles (speaker, session chair, organization staffer, etc.). During the conference people pursue their interests moving around the physical locations and engaging in different activities. A Personal Representative Agent (PRA) is an agent inhabiting the virtual space that is in charge of advancing some particular interest of a conference attendee by searching for information and talking to other software agents."

The example presented in Esteva (2003) has structured the application in six different scenes: Information Gathering Scene, Context Scene, Appointment Proposal Scene, Appointment Coordination Scene, Advertiser Scene, and Delivery Scene. However, in Esteva (2003) more details were provided for the scene Appointment Proposal, which allows us to use it as the focus on this paper.

The participants of this scene are two personal representative agents (PRA). The goal of the scene is to agree upon a set of topics for discussing during the appointment. The scene is played as follows:

1. One of the PRAs (PRA1) takes the initiative and sends an appointment proposal to the other PRA (PRA2), with a set of initial topics. This proposal has a time that defines its validity (clock). We will refer to the PRA1_x and to PRA2. Whenever the clock expires and PRA1_y has not answered, the scene moves to s5.

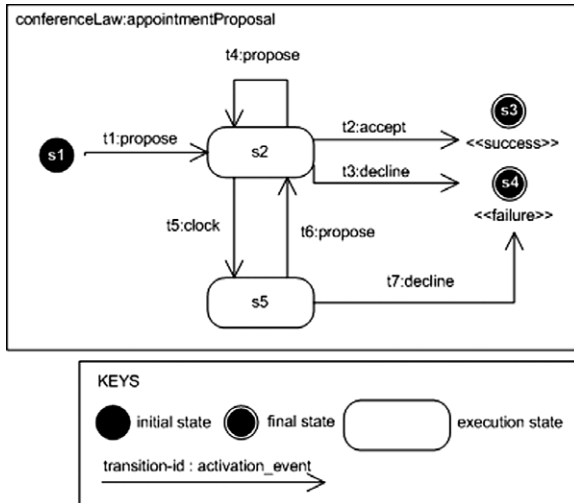


Fig. 6. Interaction protocol.

Table 16
XMLaw Code

```

01: appointmentProposal{
02:   propose($PRA1,$PRA2,$topics)
03:   accept($PRA1,$PRA2,$topics)
04:   decline($PRA1,$PRA2,$reason)
05:   s1{initial}
06:   s3{success}
07:   s4{failure}
08:   t1{s1 -> s2, propose}
09:   t2{s2 -> s3, accept}
10:   t3{s2 -> s4, decline}
11:   t4{s2 -> s2, propose}
12:   t5{s2 -> s5, clock}
13:   t6{s5 -> s2, propose}
14:   t7{s5 -> s4, decline}
15:   app-notification($PRA1, (t2),())
16:   clock{5000,regular, (t1,t4,t6),(t2,t3,t4,t5)}
17:}

```

activation is much simpler, one has only to specify which event
activates the norm.

5. Discussions

In this paper we have shown that the conceptual model of XMLaw is composed of higher-level abstractions as compared to the primitives of LGI. We have also shown that the event-based notion leads XMLaw to have a more flexible model to accommodate future changes and compose the elements when compared to EI.

To be more precise, both XMLaw and LGI deal only with the exchange of messages between agents, and are not sensitive to the internal behavior of agents, and to changes in their internal state. In general, LGI is most effective for laws that are naturally local, while XMLaw is most effective for laws that are naturally global. Laws under both approaches are not intended to specify all the details of the interaction between the agents; it is merely a constraint on the interaction. From a conceptual point of view, LGI provides a state abstraction (control state), a set of events relating to communication issues, and a set of operations for manipulating the state. The state acts basically as a hashtable where terms are stored. There is no restriction over the type of terms that can be used. This lack of a restriction may lead to a great flexibility that is useful in adapting the approach to various domains. However, a small set of high level predicates could be more useful to help in the coordination and enforcement of laws without the complexity of creating new terms. The mapping of the elements from XMLaw to LGI can be seen as creating a prolog-based model of XMLaw because the events and operations provided in LGI are general and related to low level concepts. The mapping can be used if there is a need for a decentralized architecture such as LGI. It is also important to say that although global properties can be implemented in LGI, if one uses the general solution presented in Cugola et al. (1998) and referenced in Table 4, one is not making use of the decentralized nature of LGI. On the other hand, it is also possible to write laws that make use of very specific and domain-dependent knowledge to synchronize states only when needed. However this approach introduces complexity for the specification of the laws and brings to the law specification concerns of distributed synchronization.

A flexible underlying event-based model as presented in XMLaw could make conceptual models of governance approaches more prepared to accommodate changes. We think that there is much space to improve the elements of the XMLaw model to make it more expressive and even easier for designing laws. One such improvement would be to incorporate the notion of norms described in Garcia-Camino et al. (2005).

2. PRA2 evaluates the proposal and can either (i) accept, (ii) decline, or (iii) send a counter proposal to PRA1_x with a different set of topics. The proposal has also a time that defines its validity.
3. in turn, when PRA1 receives the counter proposal of PRA2, PRA1_x evaluates this counter proposal and can also either accept, decline, or send a counter proposal to PRA2. This negotiation phase finishes when an agreement on topics is reached or one of them decides to withdraw a specific proposal.

As we have said, PRAs participate in the virtual space representing an attendee while trying to agree upon a set of topics for discussion at the appointment. Thus, when a PRA reaches an agreement for the set of topics, the PRA must inform the attendee. This is represented in XMLaw through the norm *app-notification* that can be used in other scenes to prevent agents that have not fulfilled the obligation from interacting

4.2.1. XMLaw solution

Fig. 6 shows a graphical representation based on UML state-charts of the interaction protocol of Appointment Proposal scene. The XMLaw specification of this scene is shown in Table 16. There are three types of messages: propose, accept and decline (lines 02 to 04). Those messages are used to fire most of the transitions (lines 08, 09, 10, 11, 13 and 14). The transition t5 is activated by the clock in line 16. This clock is activated every time transitions t1, t4 or t6 fire; and it is deactivated when there is a firing of transitions t2, t3, t4 or t5. This clock generates a *clock_tick* 5000 milliseconds after its activation. Line 15 declares the norm *app-notification*. This norm is given to the PRA that accepts the proposal of topics. This acceptance occurs in transition t2.

4.2.2. Discussion

When compared to the solution presented in Esteva (2003), the solution presented here has some differences: (i) the set of message definitions is reused many times in the protocol, which has lead to a much simpler protocol (for example, the number of transitions was decreased from 13 to 7); (ii) because of the event-model, the clock element is plugged into the law in order to fire transitions. When compared to EI, the transition itself has a time-out element. In other words, the transition provides the functionality of the clock. This separation leads to better separation of concerns, and better reuse once clocks can be composed with other elements; (iii) as the norm is also connected to the event-model, its

Table 17

Pros and Cons of XMLaw when compared to LGI and EI

Item	LGI	Electronic institutions	XMLaw
Level of abstraction	Low	High	High
Level of flexibility in the composition of the elements (event-based model)	High	Low	High
Scalability with decentralized laws	High	Low	Low
Global centered view of the system	No	Yes	Yes
Use of Java to help the specification of the laws	Partial (in LGI it is possible to specify the laws using a prolog-based language or pure Java, however it is not possible to combine the declarative with the imperative parts of the laws as is done with actions in XMLaw)	No	Yes
Approach is subject to a central point of congestion	No	Yes (although the enforcement is done locally, mediators have to synchronize the states to build a global view, in this way, the synchronization becomes a bottleneck)	Yes
Approach address security issues	Certification authorities	No	No
Allows agents to communicate in any different content languages	No	Partial (only prolog and lisp)	Yes
Time-sensitive norms	No	No	Yes

To summarize, XMLaw is an alternative approach to specifying laws in open multi-agent systems that presents **high level** abstractions and a flexible underlying event-based model. Thus XMLaw allows for flexible composition of the elements from its conceptual model and is flexible enough to accept new elements. **Table 17** summarizes the main contribution and limitations of the proposed model when compared to EI and LGI. The purpose of this table is to clarify when each approach should be used.

We are currently extending the XMLaw model to incorporate fault tolerance techniques. The idea is to use the laws to perform error detection and then also use laws to specify the recovery strategy through error handling (rollback, rollforward, compensation) or fault handling (diagnosis, isolation, reconfiguration, reinitialization). Thus, the XMLaw model has to evolve to accommodate the concerns related to fault handling. Other work we are performing includes using the laws to collect explicit meta-data about dependability (Kaâniche et al., 2004) using a dependability explicit computing approach. The goal is to show that our model is flexible enough to deal with very different concerns, accommodating many aspects of dependability.

6. Uncited reference

Carvalho et al. (2006a).

Acknowledgements

This work is partially supported by CNPq/Brazil under the Project “ESSMA”, Number 5520681/2002-0 and by individual grants from CNPq/Brazil.

References

- Almeida, H., Perkusich, A., Ferreira, G., Loureiro, E., Costa, E., 2006. A component model to support dynamic unanticipated software evolution. In: International Conference on Software Engineering and Knowledge Engineering (SEKE'06), 2006. San Francisco, USA. In: Proceedings of International Conference on Software Engineering and Knowledge Engineering, vol. 18, pp. 262–267.
- Arcos, J., Esteve, M., Noriega, P., Rodríguez-Aguilar, J., Sierra, C., 2005. Environment engineering for multiagent systems. Journal of Engineering Applications of Artificial Intelligence (18), 191–204.

- Ashri, R., Payne, T. R., Luck, M., Surridge, M., Sierra, C., Aguilar, J.A.R., Noriega, P., 2006. Using electronic institutions to secure grid environments. In: Proceedings of 10th International Workshop CIA on Cooperative Information Agents. Edinburgh, Scotland, pp. 461–475.
- Batista, T., Rodríguez, N., 2000. Dynamic reconfiguration of component-based applications. In: Proceedings of the International Symposium on Software Engineering for Parallel and Distributed Systems. IEEE Computer Society, pp. 32–39.
- Bou, E., López-Sánchez, M., Rodríguez-Aguilar, J., 2006. Norm adaptation of autonomic electronic institutions with multiple goals. In: ITSSA Journal International Transactions on Systems Science and Applications, vol. 1(3), pp. 227–238 (ISSN 1751-1461 (Print); ISSN 1751-147X (CD-ROM)).
- Carvalho, G., Brandão, A., Paes, R., Lucena, C., 2006. Interaction laws verification using knowledge-based reasoning. In: Workshop on Agent-Oriented Information Systems (AOIS-2006) at AAMAS.
- Carvalho, G., Lucena, C., Paes, R., Briot, J., 2006. Refinement operators to facilitate the reuse of interaction laws in open multi-agent systems. In: Proceedings of the 2006 International Workshop on Software Engineering For Large-Scale Multi-Agent Systems.
- Carvalho, G., Almeida, H., Gatti, M., Vinicius, G., Paes, R., Perkusich, A., Lucena, C., 2006. Dynamic law evolution in governance mechanisms for open multi-agent systems. In: Second Workshop on Software Engineering for Agent-oriented Systems.
- Cugola, G., Di Nitto, E., Fuggetta, A., 1998. Exploiting an event-based infrastructure to develop complex distributed systems. In: Proceedings of the 20th International Conference on Software Engineering (ICSE 98), Kyoto, Japan.
- Cuní, G., Esteve, M., García, P., Puertas, E., Sierra, C., Solchaga, T., 2004. MASFIT: multi-agent system for fish trading. In: Proceedings of the 16th European Conference on Artificial Intelligence. València, Spain, pp. 710–714.
- de C. Gatti, M.A., de Lucena, C.J., Briot, J., 2006. On fault tolerance in law-governed multi-agent systems. In: Proceedings of the 2006 International Workshop on Software Engineering For Large-Scale Multi-Agent Systems (Shanghai, China, May 22–23). SELMAS'06. ACM Press, New York, NY, 21–28. <<http://doi.acm.org/10.1145/1138063.1138068>>.
- Dignum, V., Vazquez-Salceda, J., Dignum, F., 2004. A model of almost everything: norms, structure and ontologies in agent organizations. In: Proceedings of the Third International Joint Conference on Autonomous Agents and Multiagent Systems – Volume 3 (New York, NY, July 19–23). International Conference on Autonomous Agents. IEEE Computer Society, Washington, DC, 1498–1499. <<http://dx.doi.org/10.1109/AAMAS.2004.20>>.
- Domain Language Inc., 2007. The challenge of complexity. <<http://domainrivendesign.org/>> (accessed January 2007).
- Esteve, M., 2003. Electronic institutions: from specification to development. Ph.D. Thesis. Technical University of Catalonia.
- García-Camino, A., Noriega, P., Rodríguez-Aguilar, J.A., 2005. Implementing norms in electronic institutions. In: Proceedings of the Fourth International Joint Conference on Autonomous Agents and Multiagent Systems (The Netherlands, July 25–29, 2005). AAMAS'05. ACM Press, New York, NY, pp. 667–673 <<http://doi.acm.org/10.1145/1082473.1082575>>.
- Kaâniche, M., Laprie, J.-C., Blanquart, J.-P., 2004. A dependability-explicit model for the development of computing systems. Lecture Notes in Computer Science, vol. 1943/2000.

- Lindermann, G., Ossowski, S., Padget, J., Vázquez Salceda, J., 2006. International workshop on agents, norms and institutions for regulated multiagent systems (ANIREM 2005). <<http://platon.escet.urjc.es/ANIREM2005/>> (accessed December 2006).
- Meier, R., Cahill, V., 2005. Taxonomy of distributed event-based programming systems. *The Computer Journal* 48 (5), 602–626.
- Meyer, B., 2003. The power of abstraction, reuse and simplicity: an object-oriented library for event-driven design. In: Olaf Owe et al., (Eds.), *Festschrift in Honor of Ole-Johan Dahl*. In: *Lecture Notes in Computer Science*, vol. 2635. Springer-Verlag.
- Minsky, N.H., 2003. On conditions for self-healing in distributed software systems. In: In the Proceedings of the International Autonomic Computing Workshop, Seattle, Washington.
- Minsky, N., 2005. Law governed interaction (LGI): a distributed coordination and control mechanism – An introduction and a reference manual. <<http://www.moses.rutgers.edu/documentation/manual.pdf>> (accessed 14.01.07).
- Minsky, N., 2005. On a principle underlying self-healing in heterogeneous software. *Journal of Integrated Computer-Aided Engineering*.
- Minsky, N., Murata, T., 2004. On manageability and robustness of open multi-agent systems. In: *Software Engineering for Multi-Agent Systems II. Lecture Notes in Computer Science*, pp. 189–206.
- Minsky, N.H., Rozenshtein, D., 1987. A law-based approach to object-oriented programming. In: N. Meyrowitz, (Ed.), *Conference Proceedings on Object-Oriented Programming Systems, Languages and Applications* (Orlando, Florida, United States, October 04–08, 1987). OOPSLA'87. ACM Press, New York, NY, 482–493. <<http://doi.acm.org/10.1145/38765.38851>>.
- Minsky, N.H., Ungureanu, V., 2000. Law-governed interaction: a coordination and control mechanism for heterogeneous distributed systems. *ACM Transactions on Software Engineering Methodology* 9 (3), 273–305.
- Murata, T., Minsky N., 2003. On monitoring and steering in large scale multiagent systems. In: In the Proceedings of the Second International Workshop on Large Scale Multi Agent Systems, Portland Oregon.
- Noriega, P., 1997. Agent mediated auctions: The fishmarket metaphor. Ph.D. Thesis. Universitat Autònoma de Barcelona.
- Paes, R.B., Carvalho G.R., Lucena, C.J.P., Alencar, P.S.C., Almeida H.O., Silva, V.T., 2005. Specifying laws in open multi-agent systems. In: *Agents, Norms and Institutions for Regulated Multi-agent Systems (ANIREM), AAMAS2005*.
- Paes, R., Gatti, M., Carvalho, G., Rodrigues, L., Lucena, C., 2006. A middleware for governance in open multi-agent systems. Technical Report 33/06, PUC-Rio, 14p.
- Paes, R.B., Gatti, M.A.C., Carvalho, G.R., Rodrigues, L.F.C., Lucena, C.J.P., 2006. A middleware for governance in open multi-agent systems. Technical Report 33/06, PUC-Rio, 14p.
- Paes, R., Carvalho, G., Gatti, M., Lucena, C., Briot, J., Choren, R., 2007. Enhancing the environment with a law-governed service for monitoring and enforcing behavior in open multi-agent systems. In: Weyns, D., Parunak, H.V.D., Michel, F. (Eds.), *Environments for Multi-Agent Systems, Lecture Notes in Artificial Intelligence*, vol. 4389. Springer-Verlag, Berlin, pp. 221–238.
- Rodrigues, L., Carvalho, G., Paes, R., Lucena, C., 2005. Towards an integration test architecture for open MAS. In: *Software Engineering for Agent-oriented Systems (SEAS 05)*. Uberlândia, Brasil.
- Rodriguez-Aguilar, J., 2003. On the design and construction of agent-mediated electronic institutions. Ph.D. Thesis. Volume 14 of *Monografies de l'Institut d'Investigació en Intel·ligència Artificial*. Consejo Superior de Investigaciones Científicas.
- XMLaw Specification: version 1.0, [Tech Report](#), PUC-Rio, [in press](#).

862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
Q2 888
889